H3 · 1344 × 768 · 345 frames

SINGLE CARD / BF16 vs W8A8 INT8 / MATCHED OFFLOAD

Ascend 910B3

19.66 min

Full · BF16 · 8 steps

5.64 min

Veda + LightX2V8 · 90% · INT8

3.48× vs BF16 Full8

BF16 · Offload

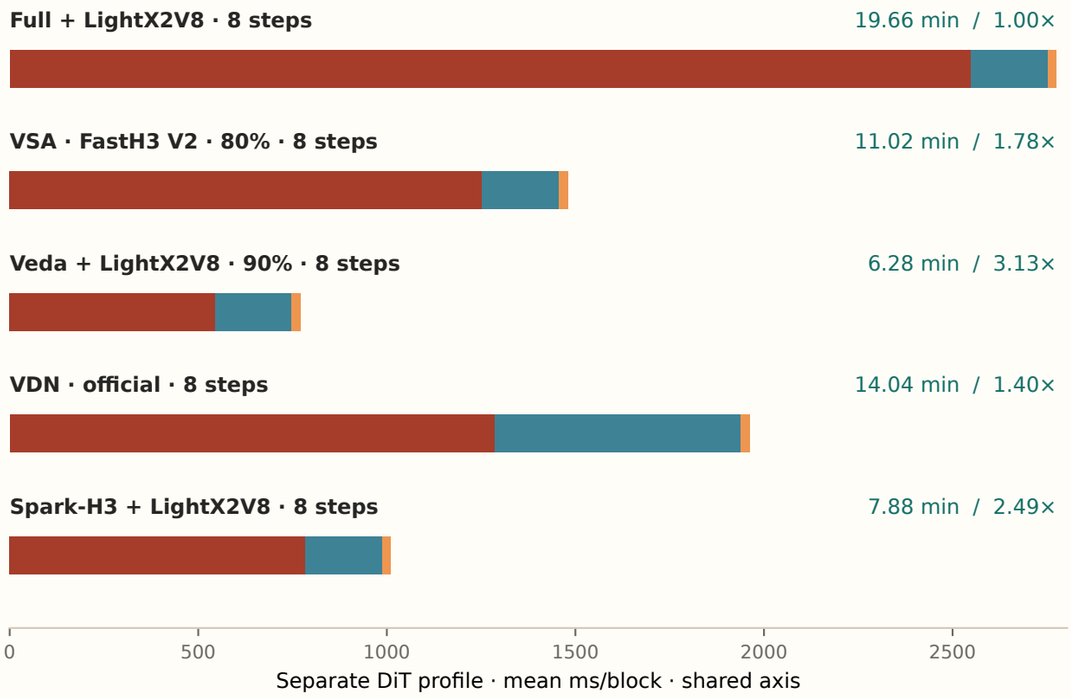

INT8 · Offload

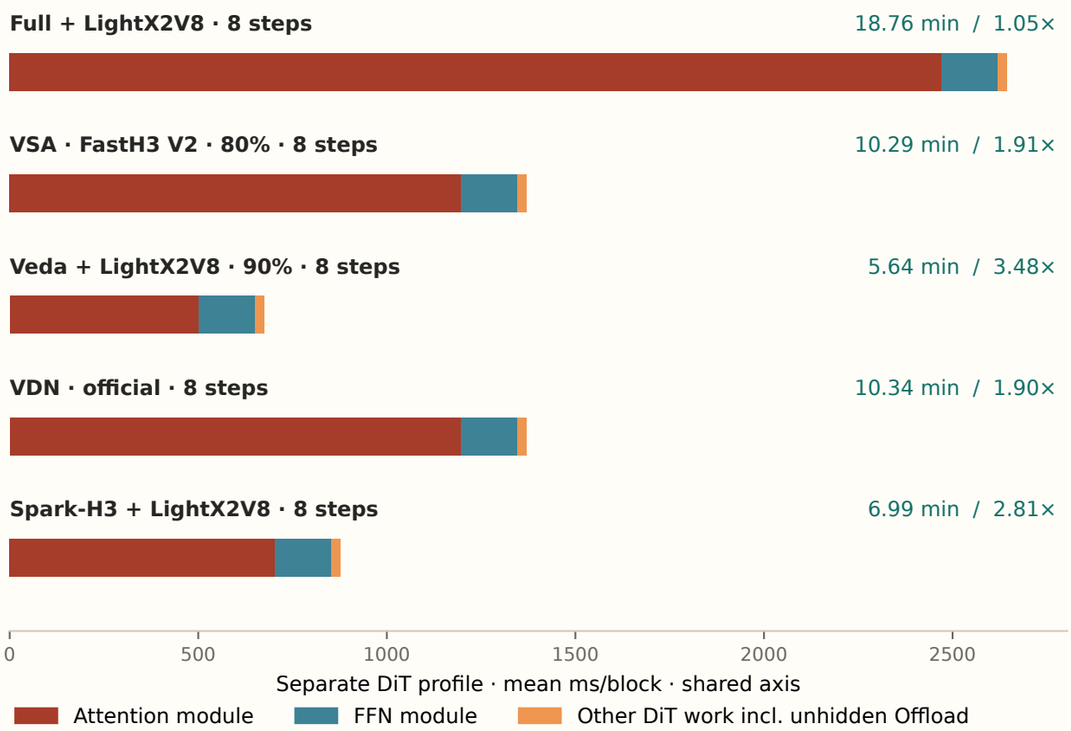

RTX PRO 6000

7.80 min

Full · BF16 · 8 steps

2.63 min

Veda + LightX2V8 · 90% · INT8

2.96× vs BF16 Full8

BF16 · Offload

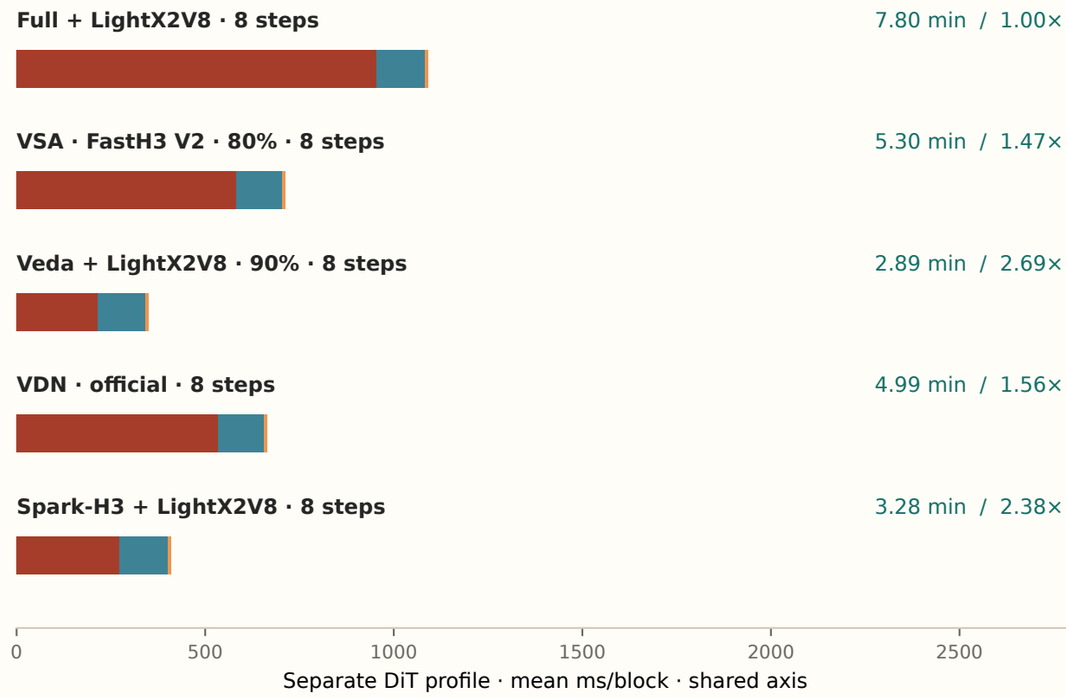

INT8 · Offload

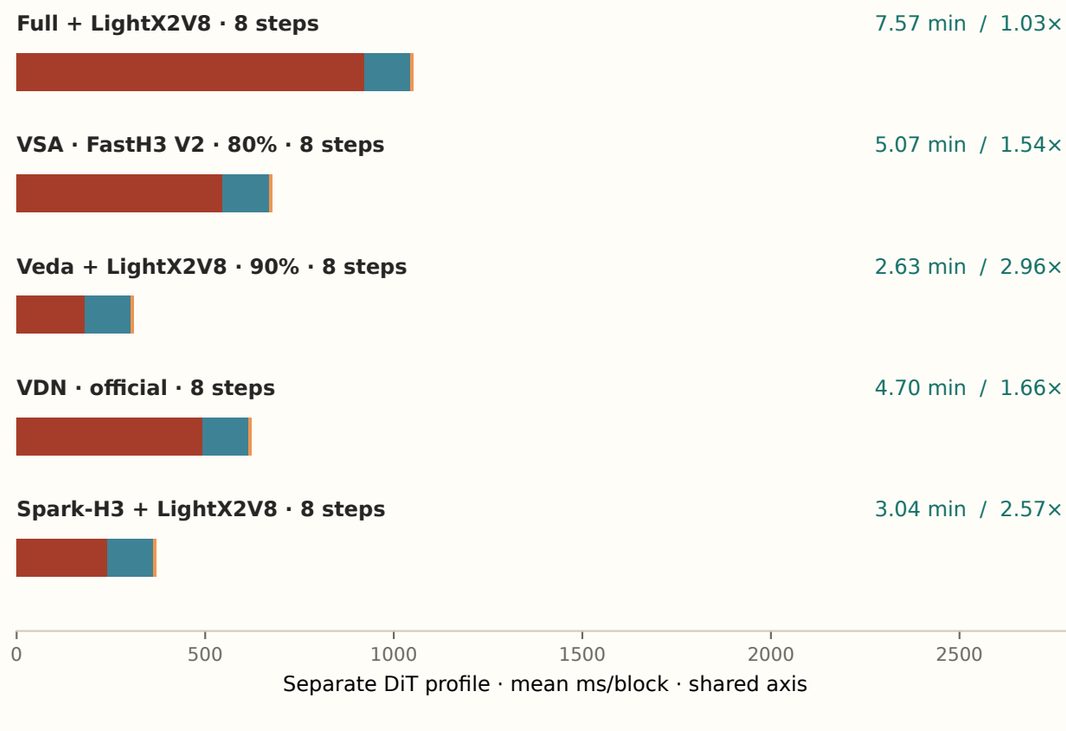

Headline times: median uninstrumented request → completed MP4; N=1 formal requests per cell. Every speedup uses same-hardware BF16 Full8, including INT8 candidates.
Bars: separate DiT profiling, normalized over actual forwards × 50 blocks; no E2E extrapolation. Offload policy: omni-layer-prefetch1-vae-resident-v1.

All methods: eight forwards, one reviewed seed, trained schedules. Module spans include waits; bars are not isolated kernel compute.